Table 3 pilot: all members' current results (week of 2026-10-08, A/B/C/D)

	Examiner Prediction			Examiner Generation			Applicant Response			Trajectory Forecasting		State Tracking		Personalized Modeling	
Model	Occ.	Type	Grnd.	R-L	MET.	Judge	Strat.	Edit	Supp.	C-idx	IBS↓	Claim	Rej.	ΔOcc.	ΔR-L
Proprietary															
gpt-5-mini	0.55 n=99	0.18 n=99	0.22 n=62	0.17 n=39/50	0.22 n=39/50	2.05 n=39/50	0.52 n=10	0.16 n=10	0.53 n=49/53	0.49 n=100	0.25 n=100	0.78 n=10	1.00 n=10	-0.11 n=14	-0.01 n=6/8
Open-weight															
Qwen3.6-27B	0.65 n=99	0.41 n=99	0.29 n=62	0.15 n=44/50	0.11 n=44/50	2.50 n=44/50	0.54 n=10	0.25 n=10	0.48 n=23/26	–	–	0.43* n=10	0.67* n=10	+0.27 n=14	+0.04 n=6/8
Input ablation															
gpt-5-mini, history only	–	–	–	–	–	–	–	–	–	0.55 n=100	0.22 n=100	–	–	–	–
Reference (no language model)															
Cox PH, cutoff features	–	–	–	–	–	–	–	–	–	0.54 n=100	0.19 n=100	–	–	–	–
Kaplan-Meier / Aalen-Johansen	–	–	–	–	–	–	–	–	–	0.50 n=100	0.19 n=100	–	–	–	–

Pilot results on small fixed subsets (data_v1_share v1.2), not benchmark findings. n = scored cases (scored/eligible when some failed). Prediction, generation, trajectory and profile cells use the structured state (A's GPT-5-mini A1 state) + history; State Tracking uses the record. Claim and Rej. are at the A1 cutoff on 10 human-reviewed gold apps; Strat., Edit, Supp. on the same 10 apps (Supp. n = verifiable argument units).
* 5 of 10 Qwen A1 snapshots were cut at 16,384 tokens and scored as empty; on the 14 snapshots both models parsed, Claim is 0.90 (Qwen) vs 0.80 (gpt-5-mini). Judge = fixed GPT-5-mini judge, 1-5.
Examiner Prediction and Personalized cells use automatic, not yet reviewed gold (B). Δ = with profile minus without on a small paired subset. Qwen uses provider-default thinking, gpt-5-mini reasoning_effort=low. – = not run.